

# Brief Introduction to Reinforcement Learning

## Examples in Q-learning

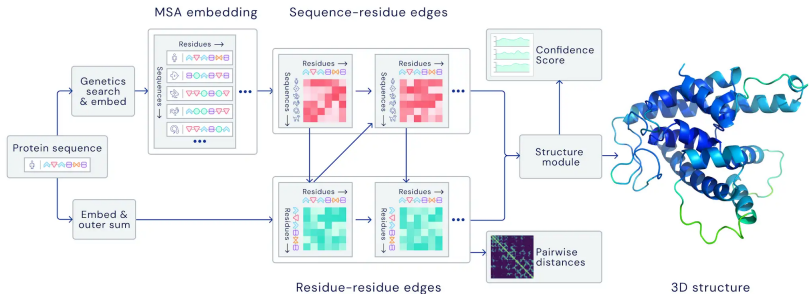
Eric Qu

Duke Kunshan University

December 15, 2020

# AlphaFold

- Recently, Google has made huge progress in protein folding.
- The program AlphaFold could shorten the time to work out certain structure of proteins given certain amino acids sequences from a few weeks of experiments to a few hours of computation.
- This is "a solution to a 50-year-old grand challenge in biology".
- The algorithm they use is - **Reinforcement Learning**



# Markov Process (Recall)

- Markov Property




$$\mathbb{P}(S_{t+1}|S_t) = \mathbb{P}(S_{t+1}|S_1, S_2, \dots, S_t)$$

- Markov Process (Markov Chain)

- Memoryless stochastic process defined by  $\langle \mathcal{S}, \mathcal{P} \rangle$
  - $\mathcal{S}$  State Space
  - $\mathcal{P}$  Transition Probability Matrix,  $\mathcal{P}_{ss'} = P(S_{t+1} = s' | S_t = s)$

- Example:



$$\mathcal{P} = \begin{bmatrix} 0 & 0.75 & 0 & 0.25 \\ 0.5 & 0 & 0.5 & 0 \\ 0 & 1 & 0 & 0 \\ 0.25 & 0.5 & 0.25 & 0 \end{bmatrix}$$

# Markov Reward Process

## Definition

- Defined by  $\langle \mathcal{S}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
- Reward  $\mathcal{R}$ 
  - $R_t$ : Reward in time  $t$ .
  - Reward in state  $s$  is: In some time  $t$  at state  $s$ , the expected reward at  $t + 1$ .

$$R_s = \mathbb{E}[R_{t+1} | S_t = s]$$

- Why  $t + 1$ ? It is defined to get a reward when leaving a state.
- Discount factor  $\gamma \in [0, 1]$
- Return
  - Return  $G_t$  of a time  $t$  is the sum of discounted reward after  $t$ .

$$G_t = R_{t+1} + \gamma R_{t+2} + \cdots = \sum_{i=0}^{\infty} \gamma^i R_{t+i+1}$$

- $\gamma$ : immediate return or long-term return

# Markov Reward Process


## Value Function

- Value Function
  - Expected return starting from state  $s$ .

$$v(s) = \mathbb{E}[G_t | S_t = s]$$

- Example

Let  $\gamma = 0.5$



$$S_1 \rightarrow S_2 \rightarrow S_1 \rightarrow S_4 \rightarrow S_3:$$
$$G_1 = -1 + 0.5(-2) + (0.5)^2(-1) + (0.5)^3 10$$

# Markov Reward Process

## Ballman Function

$$\mathbb{E}[\mathbb{E}[G_{t+1}|S_{t+1}, S_t]|S_t] = \mathbb{E}[G_{t+1}|S_t]$$

$$\begin{aligned}\mathbb{E}[\mathbb{E}[G_{t+1}|S_{t+1}, S_t]|S_t] &= \mathbb{E}[\mathbb{E}[G_{t+1}|S_{t+1}]|S_t] \\ &= \mathbb{E}[v(s+1)|S_t]\end{aligned}$$

$$\mathbb{E}[G_{t+1}|S_t] = \mathbb{E}[v(s+1)|S_t]$$

$$v(s) = \mathbb{E}[G_t|S_t = s]$$

$$\begin{aligned}&= \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \dots) | S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma G_{t+1} | S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) | S_t = s]\end{aligned}$$

- Ballman function:  $v(s) = \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) | S_t = s]$
- Two parts: the immediate reward  $R_{t+1}$ , and the discounted value of successor state.
- Also:

$$v(s) = R_s + \gamma \sum_{s' \in \mathcal{S}} P_{ss'} v(s')$$

# Markov Decision Process


## Definition

- A Markov Decision Process is defined by  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
- $\mathcal{A}$  is the Action Space. Here, each  $\mathcal{P}$  and  $\mathcal{R}$  is corresponding to an action  $a$ :

$$\mathcal{P}_{ss'}^a = P(S_{t+1} = s' | S_t = s, A_t = a)$$

$$\mathcal{R}_s^a = \mathbb{E}[\mathcal{R}_{t+1} | S_t = s, A_t = a]$$

- Example



Action Space:

$$\mathcal{A} = \{R, L, U, D, UL, UR, DL, DR\}$$

Each action is corresponding to a reward.

# Markov Decision Process

## Policy

- **Policy**

- Policy  $\pi$  describes all possible behavior of the agent.  $\pi(a|s)$  is the probability of taking action  $a$  in state  $s$ .

$$\pi(a|s) = P(A_t = a | S_t = s)$$

- Given a  $M = \langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$  and a policy  $\pi$ , then  $S_1, S_2, \dots$  is a Markov Process  $\langle \mathcal{S}, \mathcal{P}^\pi \rangle$ , with,

$$\mathcal{P}_{ss'}^\pi = \sum_{a \in \mathcal{A}} \pi(a|s) \mathcal{P}_{ss'}^a$$

$S_1, R_1, S_2, R_2, \dots$  is a Markov Decision Process  $\langle \mathcal{S}, \mathcal{P}^\pi, \mathcal{R}^\pi, \gamma \rangle$

$$\mathcal{R}_s^\pi = \sum_{a \in \mathcal{A}} \pi(a|s) \mathcal{R}_s^a$$



# Markov Decision Process

## Value Function

- **Value Function**

- **State Value Function:** The expected return starting from state  $s$ , following policy  $\pi$ .

$$v_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$$

- **State-action Value Function:** The expected return starting from state  $s$ , take action  $a$ , following policy  $\pi$ .

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a]$$

- 

$$v_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) q_{\pi}(s, a)$$

$$q_{\pi}(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_{\pi}(s')$$

# Markov Decision Process

## Bellman Expectation Function

- We could combine them:

$$v_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_{\pi}(s') \right)$$

$$q_{\pi}(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \sum_{a' \in \mathcal{A}} \pi(a'|s') q_{\pi}(s', a')$$

- **Bellman Expectation Function**

$$v_{\pi}(s) = \mathbb{E}_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s]$$

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) | S_t = s, A_t = a]$$

# Markov Decision Process

## Optimal Value Function

- **Optimal State Value Function:** The maximum state value function over all policies.

$$v_*(s) = \max_{\pi} v_{\pi}(s)$$

- **Optimal State-action Value Function:** The maximum state-action value function over all policies.

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$

- **Optimal Policy:** For every MDP, there exists an optimal policy  $\pi_*$  that is better or equal to all other policies and achieve both optimal value functions:  $v_{\pi_*}(s) = v_*(s), q_{\pi_*}(s, a) = q_*(s, a)$   
The ultimate goal of Reinforcement Learning is to find the optimal policy.

# Markov Decision Process

## Bellman Optimality Function

- We have

$$v_*(s) = \max_a q_*(s, a)$$

$$q_*(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_*(s')$$

- Combine them, then we have the **Bellman Optimality Function**

$$v_*(s) = \max_a \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_*(s') \right)$$

$$q_*(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \max_{a'} q_*(s', a')$$

# Q-learning

## Introduction

- Q-learning was a big break through in the early days of Reinforcement Learning.
- For target policy, Q-learning updates it in a greedy way:

$$\pi(s_{t+1}) = \arg \max_{a'} q(s_{t+1}, a')$$

For every step, choose the action with the max expected return.



- For  $q(s, a)$ , it will be updated by Bellman Optimality Function:

$$q(s, a) \leftarrow (1 - \alpha)q(s, a) + \alpha[\mathcal{R}_s^a + \gamma \max_{a'} q(s', a')]$$

# Q-learning


## Graph Example

- We could try to train an agent to evacuate from a room.
- The room could be represented by a graph, room as node and door as edge.



- The should first identify the rule of evacuation.
  - We start from a random room and try to get to the outside (5).
  - $\mathcal{S} = \{0, 1, 2, 3, 4, 5\}, \mathcal{A} = \{0, 1, 2, 3, 4, 5\}, \mathcal{P}_{ss'}^a = \begin{cases} 1 & s' = a \\ 0 & s' \neq a \end{cases}$
- Then, we can construct a reward matrix for the game.
- Note that the construction of reward is of great significance in reinforcement learning.

State	Action					
	0	1	2	3	4	5
0	-1	-1	-1	-1	0	-1
1	-1	-1	-1	0	-1	100
2	-1	-1	-1	0	-1	-1
3	-1	0	0	-1	0	-1
4	0	-1	-1	0	-1	100
5	-1	0	-1	-1	0	100



- Then we set  $\gamma = 0.8$  and start Q-learning algorithm.
- During the algorithm, we aim to maintain a Q-matrix that stores all  $q(s, a)$  values.
- We initialize the Q-matrix to be all zeros.
- Assume we randomly start at state 1.

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

$$R = \begin{matrix} & \begin{matrix} \text{Action} \\ 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} \text{State} \\ 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix} \end{matrix}$$



- Assume we randomly start at state 1.
- Then we apply a  $\epsilon$ -greedy algorithm to choose the next state:
  - From 1, we could reach 3 and 5.
  - $\epsilon$ -greedy algorithm is having probability  $(1 - \epsilon)$  to choose according to  $q(s, a)$  and probability  $\epsilon$  to choose a random action.
  - Purpose of  $\epsilon$ -greedy: prevent it to stuck in a local maximum.
  - Since the Q table is empty now, we should randomly choose a state from 3 and 5. Assume we choose 5.

$$Q = \begin{array}{c} \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} \end{array} \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$$R = \begin{array}{c} \begin{matrix} \text{State} \end{matrix} \begin{matrix} \text{Action} \\ 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} \end{array} \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix}$$

- Assume we randomly start at state 1 and go to state 5.
- After reaching state 5, we should update the Q table:
  - For simplicity concern, we ignore the learning rate  $\alpha$
  - $q(s, a) \leftarrow q(s, a) + \mathcal{R}_s^a + \gamma \max_{a'} q(s', a')$
  - From 5, we could go to state 1, 4 or 5. But the q-value of all these actions are 0.
  - $q(1, 5) = q(1, 5) + \mathcal{R}_1^5 + 0.8 \max[q(5, 1), q(5, 4), q(5, 5)] = 100$
- After reaching state 5, we finish one "episode".

$$Q = \begin{array}{c|cccccc} & 0 & 1 & 2 & 3 & 4 & 5 \\ \hline 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 100 \\ 2 & 0 & 0 & 0 & 0 & 0 & 0 \\ 3 & 0 & 0 & 0 & 0 & 0 & 0 \\ 4 & 0 & 0 & 0 & 0 & 0 & 0 \\ 5 & 0 & 0 & 0 & 0 & 0 & 0 \end{array}$$

$$R = \begin{array}{c|cccccc} & \text{Action} & 0 & 1 & 2 & 3 & 4 & 5 \\ \hline \text{State} & 0 & -1 & -1 & -1 & -1 & 0 & -1 \\ 1 & -1 & -1 & -1 & 0 & -1 & 100 \\ 2 & -1 & -1 & -1 & 0 & -1 & -1 \\ 3 & -1 & 0 & 0 & -1 & 0 & -1 \\ 4 & 0 & -1 & -1 & 0 & -1 & 100 \\ 5 & -1 & 0 & -1 & -1 & 0 & 100 \end{array}$$

- For the second episode, assume we start at state 3.
- From  $\mathcal{R}$ , we have 3 possible actions: go to state 1, 2 or 4.  
Suppose we choose state 1.
- At state 1, we should update the Q table:  

$$q(3, 1) = q(3, 1) + \mathcal{R}_3^1 + 0.8 \max[q(1, 5), q(1, 3)] = 80$$
- The next step should be the same as above.

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 80 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$


$$R = \begin{matrix} & \begin{matrix} \text{Action} \\ 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} \text{State} \\ 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix} \end{matrix}$$

- After a few rounds of training, the Q table converged to the left.
- We normalize it by dividing the largest value.

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 400 & 0 \\ 0 & 0 & 0 & 320 & 0 & 500 \\ 0 & 0 & 0 & 320 & 0 & 0 \\ 0 & 400 & 256 & 0 & 400 & 0 \\ 320 & 0 & 0 & 320 & 0 & 500 \\ 0 & 400 & 0 & 0 & 400 & 500 \end{bmatrix} \end{matrix}$$

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 80 & 0 \\ 0 & 0 & 0 & 64 & 0 & 100 \\ 0 & 0 & 0 & 64 & 0 & 0 \\ 0 & 80 & 51 & 0 & 80 & 0 \\ 64 & 0 & 0 & 64 & 0 & 100 \\ 0 & 80 & 0 & 0 & 80 & 100 \end{bmatrix} \end{matrix}$$


- The final graph looks like this.
- The optimal path is just following the max value.



$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 80 & 0 \\ 0 & 0 & 0 & 64 & 0 & 100 \\ 0 & 0 & 0 & 64 & 0 & 0 \\ 0 & 80 & 51 & 0 & 80 & 0 \\ 64 & 0 & 0 & 64 & 0 & 100 \\ 0 & 80 & 0 & 0 & 80 & 100 \end{bmatrix} \end{matrix}$$

# Flappy Bird Example

- Let look at a more interesting example.
- [Flappy Bird](#)



- State  $\mathcal{S}$


- Obviously, we could choose each frame of game to be the state.
- For simplicity concern, we choose  $(\Delta x, \Delta y)$  to be the state.
- Since we could not make the Q-table infinite, the  $(\Delta x, \Delta y)$  are floored according to pixels.

- Action  $\mathcal{A}$

- In each frame, it could fly up or do nothing.

- Reward  $\mathcal{R}$

- For each frame, when the bird is alive, we give it reward 1. When the bird dies, we give it reward -1000.



- We try to train a Q-table, such as this:

State	Fly	Don't Fly
$(\Delta x_1, \Delta y_1)$	3	20
$(\Delta x_1, \Delta y_2)$	100	-20
...	...	...
$(\Delta x_m, \Delta y_{n-1})$	-100	1
$(\Delta x_m, \Delta y_n)$	3	-400

- The total number of states are  $n \times m$ , and each state has two actions.
- Theoretically, we could choose the maximum value each frame and survive infinite times.



- We could should the Q-learning algorithm more vigorously:

---

**Algorithm 1: Q-learning**

---

```
1 Initialization  $Q = \{ \}$ ;  
2 while  $Q$  not converge do  
3   Initialize state  $s$ , start a new game;  
4   while  $s$  is not dead state do  
5      $a = \pi(s)$ ; // Use  $\epsilon$ -greedy.  
6     Use  $a$  in game, get new state  $s'$  and reward  $\mathcal{R}_s^a$   
7      $q(s, a) \leftarrow (1 - \alpha)q(s, a) + \alpha[\mathcal{R}_s^a + \gamma \max_{a'} q(s', a')]$   
8      $S \leftarrow S'$   
9   end  
10 end
```

---

- Let's try it !
- Flappy Bird

# Thank You!